

FedIRT: An R package and shiny app for estimating federated item response theory models

Biying Zhou¹ and Feng Ji¹

¹ Department of Applied Psychology & Human Development, University of Toronto, Toronto, Canada
Corresponding author

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#)
- [Repository](#)
- [Archive](#)

Editor: [Open Journals](#)

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

We developed an R package, FedIRT, to estimate item response theory (IRT) models—including 1PL, 2PL, and graded response models—with additional privacy features. This package enables parameter estimation in a distributed manner without compromising accuracy, leveraging recent advances in federated learning. Numerical experiments demonstrate that federated IRT estimation achieves statistical performance comparable to mainstream IRT packages in R, with the added benefits of privacy preservation and minimal communication costs. The R package also includes a user-friendly Shiny app that allows clients (e.g., individual schools) and servers (e.g., school boards) to easily apply our proposed method.

Statement of Need

IRT ([Embretson & Reise, 2013](#)) is a statistical modeling framework grounded in modern test theory, frequently used in the educational, social, and behavioral sciences to measure latent constructs through multivariate human responses. Traditional IRT estimation mandates the centralization of all individual raw response data in one location, which potentially compromises the privacy of the data and participants ([Lemons, 2014](#)).

Federated learning has emerged as a field addressing data privacy issues and techniques for parameter estimation in a decentralized, distributed manner. However, there is currently no package available in psychometrics, especially in the context of IRT, that integrates federated learning with IRT model estimation.

Popular IRT packages in R, such as `mirt` ([Chalmers, 2012](#)) and `ltm` ([Rizopoulos, 2007](#)), require storing and computing all data in a single location, which can potentially lead to violations of privacy policies when dealing with highly sensitive data (e.g., high-stakes student assessment data).

Therefore, we have developed a specialized R package, FedIRT, which integrates federated learning with IRT and includes an accompanying Shiny app designed to address real-world implementation challenges and reduce the burden of learning R programming for users. This app implements the method in a user-friendly and accessible manner.

Method

Here we briefly introduce the key idea behind integrating federated learning with IRT. For technical details, please refer to our methodological discussions on Federated IRT ([Zhou & Ji, 2023, 2024, In submission](#)).

Model formulation

The two-parameter logistic (2PL) IRT model is often considered the most popular IRT model in practice. In the 2PL model, the response of person i to item j is binary ($X_{ij} \in 0, 1$), and the probability that person i answers item j correctly, given discrimination parameter α_j and difficulty parameter β_j , is given by:

$$P(X_{ij} = 1|\theta_i) = \frac{e^{\alpha_j(\theta_i - \beta_j)}}{1 + e^{\alpha_j(\theta_i - \beta_j)}}$$

To make our package available for polytomous response, we also developed a federated learning estimation algorithm for the Generalized Partial Credit Model (GPCM) in which the probability of a person with the ability θ_i obtaining x scores in item j is:

$$P^{\text{GPCM}}(X_{ij} = x|\theta_i) = \frac{e^{\sum_{h=1}^x \alpha_j(\theta_i - \beta_{jh})}}{\sum_{c=0}^{m_j} e^{\sum_{h=1}^c \alpha_j(\theta_i - \beta_{jh})}}$$

In this function, β_{jh} is the difficulty of scoring level h for item j , and for each item j , all difficulty levels have the same discrimination α_j . m_j is the maximum score of item j .

Model estimation

In both the 2PL and GPCM models, we often assume that ability follows a standard normal distribution, allowing us to apply marginal maximum likelihood estimation (MMLE).

We use a combination of traditional MMLE with federated average (FedAvg) and federated stochastic gradient descent (FedSGD) (McMahan et al., 2017). In our case, the log-likelihood and partial gradients are sent from the clients to the server. The server then uses FedSGD to update the item parameters and sends them back to the clients.

Taking the 2PL model as an example, the marginal log-likelihood function l for each school k can be approximated using Gaussian-Hermite quadrature with q equally-spaced levels. Let $V(n)$ be the ability value at level n , and $A(n)$ be the weight at level n .

$$l_k \approx \sum_{i=1}^{N_k} \sum_{j=1}^J X_{ijk} \times \log\left[\sum_{n=1}^q P_j(V(n))A(n)\right] + (1 - X_{ijk}) \times \log\left[\sum_{n=1}^q Q_j(V(n))A(n)\right]$$

By applying FedAvg, the server collects the log-likelihood values from all k schools and then sums up all the likelihood values to get the overall log-likelihood value: $l = \sum_{k=1}^K l_k$.

The server collects a log-likelihood value l_k and all derivatives $\frac{l_k}{\partial \alpha_j}$ and $\frac{l_k}{\partial \beta_j}$ from all clients, then observe that $\frac{\partial l}{\partial \alpha_j} = \sum_{k=1}^K \frac{l_k}{\partial \alpha_j}$ and $\frac{\partial l}{\partial \beta_j} = \sum_{k=1}^K \frac{l_k}{\partial \beta_j}$ by FedSGD, the server sums up all log-likelihood values and derivative values.

Also, we provided an alternative solution, Federated Median, which uses the median of the likelihood values to replace the sum of likelihood values in Fed-MLE (Liu et al., 2020), with additional robustness to handle outliers in input data.

With estimates of α_j and β_j in 2PL or β_{jh} in GPCM, we can obtain empirical Bayesian estimates of students' ability (Bock & Aitkin, 1981).

67 Comparison with existing packages

68 We demonstrate that our package generates comparable results to established IRT packages,
69 such as mirt (Chalmers, 2012).

70 Figure 1 and Figure 2 show the comparison of the discrimination and difficulty parameters
71 between mirt and FedIRT based on example_data_2PL in our package.




Figure 1: Discrimination parameter estimates comparison




Figure 2: Difficulty parameter estimates comparison

72 Availability

73 The R package FedIRT is publicly available on [Github](#). It could be installed and run by using
74 the following commands:

```
devtools::install_github("Feng-Ji-Lab/FedIRT")
library(FedIRT)
```

Example of the integrated function

We provide a function `fedirt_file()` in the package, and the detailed usage of the function is shown in the user manual. We demonstrate an example here.

Suppose we have a dataset called `dataset.csv`, and the head of this dataset is shown below. There should be one column indicating the school, for example, "site" here. Each other column indicates an item, and each row represents an answering status.

site	X1	X2	X3	X4	X5
10	1	0	0	0	0
7	0	0	1	0	0
9	0	0	1	1	1
1	1	0	1	1	1
2	1	0	0	0	0

First, we need to read the dataset.

```
# read dataset
data <- read.csv("dataset.csv", header = TRUE)
```

Then, we call the function `FedIRT::fedirt_file()` to obtain the result. It returns a list of parameter estimates for item discriminations, item difficulties, and each sites' effect and each students' abilities.

```
# call the fedirt_file function
result <- fedirt_file(data, model_name = "2PL")
```

Finally, we can extract the results or use the parameter estimates for further analysis.

```
result$a
result$b
```

Apart from using the results for further analysis, we can also use `summary()` to generate a snapshot of the result. Here is an example below.

```
summary(result)
```

Then, the result will be printed in the console as follows:

```
Summary of FedIRT Results:
```

```
Counts:
```

```
function gradient
      735      249
```

```
Convergence Status (convergence):
```

```
Converged
```

```
Log Likelihood (loglik):
```

```
[1] -7068.258
```

```
Difficulty Parameters (b):
```

```
[1] -185.88151839  0.99524035  0.92927254  ...
```

```
Discrimination Parameters (a):
```

```
[1] 0.0028497700 0.8440140746 -0.1190176844 ...
```

```

108 Ability Estimates:
109 School 1:
110 [1] -1.127097195 -0.922572829 -0.993953038 ...
111 School 2:
112 [1] -1.41454573 1.78068772 1.87469389 ...
113 ...
114
115 End of Summary

```

116 Example of the personscore function

117 We provide a function personscore in the package to obtain ability estimates. The detailed
 118 usage of the function is shown in the user manual. We demonstrate an example here.

```

personscoreResult = personscore(result)
summary(personscoreResult)

```

119 Summary of the person score is shown below.

120 Summary of FedIRT Person Score Results:

```

121
122 Ability Estimates:
123 School 1:
124 [1] -1.127097195 -0.922572829 -0.993953038 ...
125 School 2:
126 [1] -1.41454573 1.78068772 1.87469389 ...
127 ...
128
129 End of Summary

```

130 Example of the personfit function

131 We provide a function personfit in the package. The detailed usage of the function is shown
 132 in the user manual. We demonstrate an example here.

```

personfitResult = personfit(result)
summary(personfitResult)

```

133 After getting the result, use personfit function to get the person score result from result by
 134 personfit(result).

135 Summary of FedIRT Person Fit Results:

```

136
137 Fit Estimates:
138 School 1:
139
140      Lz      Zh      Infit      Outfit
141 4  0.7584470759 0.923163304 0.002323484 0.1482672
142 16 -0.7562447025 -1.131668935 0.005457117 0.1799583
143 27  0.3417488360 0.357870094 0.005966933 0.1734402
144 33 -0.9244005411 -1.359789298 0.179834037 0.2266634
145 ...
146 School 2:
147
148      Lz      Zh      Infit      Outfit
149 5  -0.90114567 -1.175767350 0.0009824580 0.1535794
150 8  -1.47957351 -1.888763364 0.1491518127 0.2255230
151 18 -0.13292541 -0.228824721 0.1104556086 0.2007658
152 19 -0.17257549 -0.277699184 0.0075031313 0.1350857
153 ...

```

152 Standard error (SE) calculation

153 To obtain SE, we can call the `SE()` function and input a `fedirt` object to display standard
154 errors of item parameter estimates.

```
SE(result)
```

155 Below is the result of SE.

```
156 $a
157 [1] 0.0041815497 0.1638884452 0.1204696925 ...
158 $b
159 [1] 272.43863961 0.20737386 1.25896302 ...
```

160 Example of the Shiny App

161 To provide wider access for practitioners in real-world applications, we include the Shiny user
162 interface in our package. A detailed manual was provided in the package. Taking the 2PL as
163 an example, we illustrate how to use the Shiny app below.

164 In the first step, the server end (e.g., test administrator, school board) can be launched by running
165 the Shiny app `runserver()` and the client-end Shiny app can be initialized with `runclient()`
166 with the interface shown below:

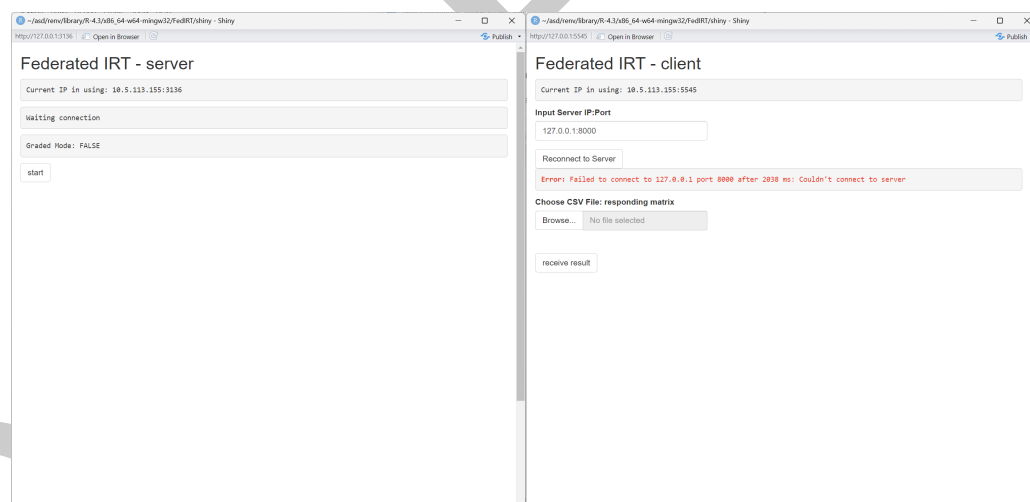


Figure 3: The initial server and client interface.

167 When the client first launches, it will automatically connect to the localhost port 8000 by
168 default.

169 If the server is deployed on another computer, type the server's IP address and port (which
170 will be displayed on the server's interface), then click "Reconnect". The screenshots of the
171 user interface are shown below.




Figure 4: Server and client interface when one school is connected.

172 Then, the client should choose a file to upload to the local Shiny app to perform local
173 calculations, without sending it to the server. The file should be a CSV file with either binary or
174 graded responses. All clients should share the same number of items and the same maximum
175 score for each item (if the responses are polytomous); otherwise, an error message will suggest
176 checking the datasets of all clients.




Figure 5: Server interface when one school uploaded dataset and client interface when a dataset is uploaded successfully.

177 After all the clients upload their data, the server should click “Start” to begin the federated
178 estimation process. After the model converges, the clients should click “Receive Result”. The
179 server will display all item parameters, and the clients will display all item parameters and
180 individual ability estimates.

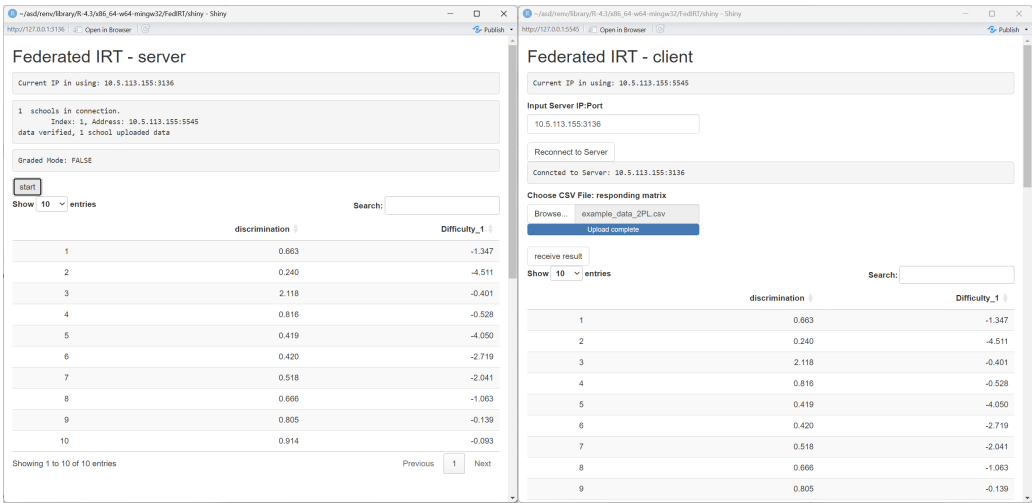


Figure 6: Server interface when estimation is completed and client interface when the results received.

181 The clients will also display bar plots of the ability estimates.

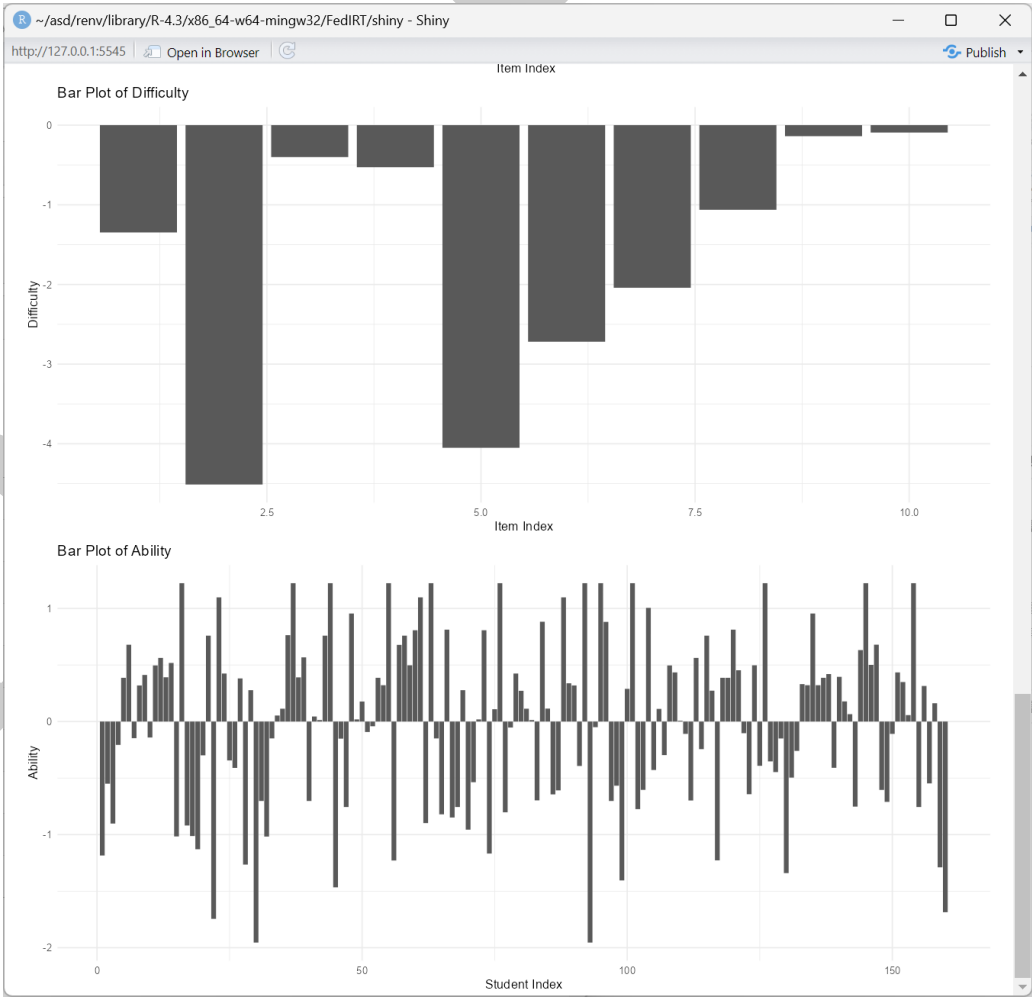


Figure 7: Client interface for displaying results.

References

- Bock, R. D., & Aitkin, M. (1981). Marginal maximum likelihood estimation of item parameters: Application of an EM algorithm. *Psychometrika*, 46(4), 443–459. <https://doi.org/10.1007/BF02293801>
- Chalmers, R. P. (2012). Mirt: A multidimensional item response theory package for the r environment. *Journal of Statistical Software*, 48, 1–29. <https://doi.org/10.18637/jss.v048.i06>
- Embretson, S. E., & Reise, S. P. (2013). *Item response theory*. Psychology Press. <https://doi.org/10.4324/9781410605269>
- Lemons, M. Q. (2014). *Predictive modeling of uniform differential item functioning preservation likelihoods after applying disclosure avoidance techniques to protect privacy* [PhD thesis]. Virginia Polytechnic Institute; State University.
- Liu, Y., Zhang, L., Ge, N., & Li, G. (2020). A systematic literature review on federated learning: From a model quality perspective. *arXiv Preprint arXiv:2012.01973*. <https://doi.org/10.48550/arXiv.2012.01973>
- McMahan, B., Moore, E., Ramage, D., Hampson, S., & Arcas, B. A. y. (2017). Communication-efficient learning of deep networks from decentralized data. *Artificial Intelligence and Statistics*, 1273–1282.
- Rizopoulos, D. (2007). Ltm: An r package for latent variable modeling and item response analysis. *Journal of Statistical Software*, 17, 1–25. <https://doi.org/10.18637/jss.v017.i05>
- Zhou, B., & Ji, F. (2023). Federated psychometrics: A distributed, privacy-preserving, and efficient IRT estimation algorithm. *APHD Research Gala*.
- Zhou, B., & Ji, F. (2024). Federated item response theory: A distributed, privacy-preserving, and efficient IRT estimation algorithm. *DPE MED Research Practicum Poster*.
- Zhou, B., & Ji, F. (In submission). *Federated item response models*.